

Overview

We conclude the course by considering hash tables, a data structure that achieves constant-time performance for core symbol table operations, provided that search keys are standard data types or simply defined. Then we consider several fundamental (and useful) examples of symbol-table clients.

Lecture: Hash Tables. We begin by describing the desirable properties of hash function and how to implement them in Java, including a fundamental tenet known as the *uniform hashing assumption* that underlies the potential success of a hashing application. Then, we consider two strategies for implementing hash tables—*separate chaining* and *linear probing*. Both strategies yield constant-time performance for search and insert under the uniform hashing assumption. We conclude with applications of symbol tables including sets, dictionary clients, indexing clients, and sparse vectors.

Exercises. *(Sorry, we are still waiting for Coursera to migrate the exercises from the old platform.)* Drill exercises on the lecture material.

Final Exam. *(Sorry, we are still waiting for Coursera to migrate the final exam from the old platform.)* The final exam is cumulative and designed to make sure you understand how each algorithm works and when it is effective. The final will not involve Java programming.

Job Interview Questions. Algorithmic interview questions based on the lecture material.

Suggested Readings. Section 3.4 in *Algorithms, 4th edition*.

✓ Complete